

Support Vector Machine

Chandan Gowda Ashwath

9/9/2020

1. Introduction:

This assignment will look at a data “juice.csv” which contains purchase information for Citrus Hill or Minute Maid orange juice.

This assignment is interested in predicting whether the customer purchased Citrus Hill or Minute Maid Orange Juice. Support Vector Machines will be used for this prediction.

We will be using three kernels i.e., linear, radial, and polynomial.

2. Data and Analysis:

Let us load the data and perform some exploratory analysis.

```
if(!require("pacman"))  
install.packages("pacman")
```

```
## Loading required package: pacman
```

```
## Warning: package 'pacman' was built under R version 3.6.3
```

```
pacman::p_load(e1071, ggplot2, caret, rmarkdown, corrplot, knitr)  
search()
```

```
## [1] ".GlobalEnv"      "package:knitr"    "package:corrplot"  
## [4] "package:rmarkdown" "package:caret"    "package:lattice"  
## [7] "package:ggplot2"  "package:e1071"    "package:pacman"  
## [10] "package:stats"    "package:graphics" "package:grDevices"  
## [13] "package:utils"    "package:datasets" "package:methods"  
## [16] "Autoloads"        "package:base"
```

```
theme_set(theme_classic())  
options(digits = 3)  
data<-read.csv("juice.csv")
```

Let us take a look at the summary and Structure of the “JUICES” data.

From the summary function, we can see that some of the variables have 5 number summaries in the form of “Min, 1st Qu, Median, Mean, 3rd Qu, Max”, when these variables should be more categorical in nature. These are not so meaningful for categorical variables.

StoreID, SpecialCH, SpecialMM and Store are examples of the categorical variables with summaries.

```
summary(data)
```

```
## Purchase WeekofPurchase StoreID PriceCH PriceMM
## CH:610 Min. :227 Min. :1.00 Min. :1.69 Min. :1.69
## MM:390 1st Qu.:239 1st Qu.:2.00 1st Qu.:1.79 1st Qu.:1.99
## Median :256 Median :3.00 Median :1.86 Median :2.09
## Mean :254 Mean :3.98 Mean :1.87 Mean :2.08
## 3rd Qu.:268 3rd Qu.:7.00 3rd Qu.:1.99 3rd Qu.:2.18
## Max. :278 Max. :7.00 Max. :2.09 Max. :2.29
## DiscCH DiscMM SpecialCH SpecialMM LoyalCH
## Min. :0.00 Min. :0.000 Min. :0.000 Min. :0.000 Min. :0.000
## 1st Qu.:0.00 1st Qu.:0.000 1st Qu.:0.000 1st Qu.:0.000 1st Qu.:0.320
## Median :0.00 Median :0.000 Median :0.000 Median :0.000 Median :0.591
## Mean :0.05 Mean :0.124 Mean :0.142 Mean :0.163 Mean :0.562
## 3rd Qu.:0.00 3rd Qu.:0.240 3rd Qu.:0.000 3rd Qu.:0.000 3rd Qu.:0.844
## Max. :0.50 Max. :0.800 Max. :1.000 Max. :1.000 Max. :1.000
## SalePriceMM SalePriceCH PriceDiff Store7 PctDiscMM
## Min. :1.19 Min. :1.39 Min. : -0.670 No :665 Min. :0.000
## 1st Qu.:1.69 1st Qu.:1.75 1st Qu.: 0.000 Yes:335 1st Qu.:0.000
## Median :2.09 Median :1.86 Median : 0.230 Median :0.000
## Mean :1.96 Mean :1.82 Mean : 0.143 Mean :0.060
## 3rd Qu.:2.13 3rd Qu.:1.89 3rd Qu.: 0.300 3rd Qu.:0.113
## Max. :2.29 Max. :2.09 Max. : 0.640 Max. :0.402
## PctDiscCH ListPriceDiff STORE
## Min. :0.0000 Min. :0.000 Min. :0.00
## 1st Qu.:0.0000 1st Qu.:0.140 1st Qu.:0.00
## Median :0.0000 Median :0.240 Median :2.00
## Mean :0.0263 Mean :0.217 Mean :1.63
## 3rd Qu.:0.0000 3rd Qu.:0.300 3rd Qu.:3.00
## Max. :0.2527 Max. :0.440 Max. :4.00
```

```
str(data)
```

```
## 'data.frame': 1000 obs. of 18 variables:
## $ Purchase : Factor w/ 2 levels "CH","MM": 2 1 2 1 1 2 1 1 1 1 ...
## $ WeekofPurchase: int 237 258 242 271 276 240 248 270 266 274 ...
## $ StoreID : int 2 7 3 2 2 1 3 1 2 7 ...
## $ PriceCH : num 1.75 1.86 1.99 1.86 1.99 1.75 1.99 1.86 1.86 1.86 ...
## $ PriceMM : num 1.99 2.18 2.23 2.18 2.18 1.99 2.23 2.18 2.18 2.13 ...
## $ DiscCH : num 0 0 0 0 0 0 0 0 0 0.47 ...
## $ DiscMM : num 0 0 0 0.06 0 0.3 0 0 0 0.54 ...
## $ SpecialCH : int 0 0 0 0 0 0 0 0 0 1 ...
## $ SpecialMM : int 0 0 0 0 1 1 0 0 0 0 ...
## $ LoyalCH : num 0.4 0.90814 0.00721 0.78839 0.97251 ...
## $ SalePriceMM : num 1.99 2.18 2.23 2.12 2.18 1.69 2.23 2.18 2.18 1.59 ...
## $ SalePriceCH : num 1.75 1.86 1.99 1.86 1.99 1.75 1.99 1.86 1.86 1.39 ...
## $ PriceDiff : num 0.24 0.32 0.24 0.26 0.19 -0.06 0.24 0.32 0.32 0.2 ...
## $ Store7 : Factor w/ 2 levels "No","Yes": 1 2 1 1 1 1 1 1 1 2 ...
## $ PctDiscMM : num 0 0 0 0.0275 0 ...
## $ PctDiscCH : num 0 0 0 0 0 ...
## $ ListPriceDiff : num 0.24 0.32 0.24 0.32 0.19 0.24 0.24 0.32 0.32 0.27 ...
## $ STORE : int 2 0 3 2 2 1 3 1 2 0 ...
```

So, let us first convert these variables to categories.

From the summary data, There are three store variables i.e, StoreID, Store7 and STORE and all these variables are related. StoreID and STORE contains same number of observations except for store 7 values converted as 0.

Also, the SpecialCH and SpecialMM show that most purchases did not include a special value on either juice brand.

```
data1<-data
cvar = c("StoreID","SpecialCH","SpecialMM","STORE")
data1[cvar] = lapply(data1[cvar], as.factor)
summary(data1)
```

```
## Purchase WeekofPurchase StoreID PriceCH PriceMM DiscCH
## CH:610 Min. :227 1:146 Min. :1.69 Min. :1.69 Min. :0.00
## MM:390 1st Qu.:239 2:204 1st Qu.:1.79 1st Qu.:1.99 1st Qu.:0.00
## Median :256 3:182 Median :1.86 Median :2.09 Median :0.00
## Mean :254 4:133 Mean :1.87 Mean :2.08 Mean :0.05
## 3rd Qu.:268 7:335 3rd Qu.:1.99 3rd Qu.:2.18 3rd Qu.:0.00
## Max. :278 Max. :2.09 Max. :2.29 Max. :0.50
## DiscMM SpecialCH SpecialMM LoyalCH SalePriceMM
## Min. :0.000 0:858 0:837 Min. :0.000 Min. :1.19
## 1st Qu.:0.000 1:142 1:163 1st Qu.:0.320 1st Qu.:1.69
## Median :0.000 Median :0.591 Median :2.09
## Mean :0.124 Mean :0.562 Mean :1.96
## 3rd Qu.:0.240 3rd Qu.:0.844 3rd Qu.:2.13
## Max. :0.800 Max. :1.000 Max. :2.29
## SalePriceCH PriceDiff Store7 PctDiscMM PctDiscCH
## Min. :1.39 Min. : -0.670 No :665 Min. :0.000 Min. :0.0000
## 1st Qu.:1.75 1st Qu.: 0.000 Yes:335 1st Qu.:0.000 1st Qu.:0.0000
## Median :1.86 Median : 0.230 Median :0.000 Median :0.0000
## Mean :1.82 Mean : 0.143 Mean :0.060 Mean :0.0263
## 3rd Qu.:1.89 3rd Qu.: 0.300 3rd Qu.:0.113 3rd Qu.:0.0000
## Max. :2.09 Max. : 0.640 Max. :0.402 Max. :0.2527
## ListPriceDiff STORE
## Min. :0.000 0:335
## 1st Qu.:0.140 1:146
## Median :0.240 2:204
## Mean :0.217 3:182
## 3rd Qu.:0.300 4:133
## Max. :0.440
```

```
str(data1)
```

```
## 'data.frame': 1000 obs. of 18 variables:
## $ Purchase : Factor w/ 2 levels "CH","MM": 2 1 2 1 1 2 1 1 1 1 ...
## $ WeekofPurchase: int 237 258 242 271 276 240 248 270 266 274 ...
## $ StoreID : Factor w/ 5 levels "1","2","3","4",...: 2 5 3 2 2 1 3 1 2 5 ...
## $ PriceCH : num 1.75 1.86 1.99 1.86 1.99 1.75 1.99 1.86 1.86 1.86 ...
## $ PriceMM : num 1.99 2.18 2.23 2.18 2.18 1.99 2.23 2.18 2.18 2.13 ...
## $ DiscCH : num 0 0 0 0 0 0 0 0 0 0.47 ...
## $ DiscMM : num 0 0 0 0.06 0 0.3 0 0 0 0.54 ...
## $ SpecialCH : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 2 ...
```

```
## $ SpecialMM      : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 1 1 ...
## $ LoyalCH        : num  0.4 0.90814 0.00721 0.78839 0.97251 ...
## $ SalePriceMM     : num  1.99 2.18 2.23 2.12 2.18 1.69 2.23 2.18 2.18 1.59 ...
## $ SalePriceCH     : num  1.75 1.86 1.99 1.86 1.99 1.75 1.99 1.86 1.86 1.39 ...
## $ PriceDiff       : num  0.24 0.32 0.24 0.26 0.19 -0.06 0.24 0.32 0.32 0.2 ...
## $ Store7          : Factor w/ 2 levels "No","Yes": 1 2 1 1 1 1 1 1 1 2 ...
## $ PctDiscMM       : num  0 0 0 0.0275 0 ...
## $ PctDiscCH       : num  0 0 0 0 0 ...
## $ ListPriceDiff    : num  0.24 0.32 0.24 0.32 0.19 0.24 0.24 0.32 0.32 0.27 ...
## $ STORE            : Factor w/ 5 levels "0","1","2","3",...: 3 1 4 3 3 2 4 2 3 1 ...
```

Let us see the correlation among the numerical variables

One can see almost perfect correlation of around 0.99 between two pairs of variables i.e., PctDiscMM and DiscMM, and PctDiscCH and DiscCH. The two pairs of variables show almost a straight line. These two pairs also have an overall correlation value of around 0.99.


There are also a couple of other highly correlated pairs with > 0.8 in magnitude as well, such as PctDicMM and PriceDiff. These show a general linear trend between the two variables that is easy to distinguish.

```
data2 <- data1[, -c(1, 3, 8, 9, 14, 18)]
cor(data2)
```

```
##           WeekofPurchase PriceCH PriceMM DiscCH DiscMM LoyalCH
## WeekofPurchase      1.00000  0.7143  0.58339  0.3746  0.23744  0.18577
## PriceCH              0.71431  1.0000  0.62382  0.1707  0.11347  0.07155
## PriceMM              0.58339  0.6238  1.00000  0.0732  0.00592  0.11031
## DiscCH               0.37464  0.1707  0.07315  1.0000  0.01902  0.13196
## DiscMM               0.23744  0.1135  0.00592  0.0190  1.00000 -0.00951
## LoyalCH              0.18577  0.0716  0.11031  0.1320 -0.00951  1.00000
## SalePriceMM          0.11236  0.2390  0.53219  0.0232 -0.84346  0.06731
## SalePriceCH          0.21173  0.5871  0.39361 -0.6974  0.06691 -0.05638
## PriceDiff            -0.00524 -0.0825  0.29334  0.3860 -0.82565  0.09254
## PctDiscMM            0.21843  0.0965 -0.01411  0.0152  0.99880 -0.01171
## PctDiscCH            0.36410  0.1532  0.06833  0.9990  0.01953  0.13131
## ListPriceDiff        0.05629 -0.1646  0.66821 -0.0702 -0.10055  0.07110
##           SalePriceMM SalePriceCH PriceDiff PctDiscMM PctDiscCH
## WeekofPurchase      0.1124    0.2117 -0.00524    0.2184    0.3641
## PriceCH              0.2390    0.5871 -0.08251    0.0965    0.1532
## PriceMM              0.5322    0.3936  0.29334   -0.0141    0.0683
## DiscCH               0.0232   -0.6974  0.38599    0.0152    0.9990
## DiscMM              -0.8435    0.0669 -0.82565    0.9988    0.0195
## LoyalCH              0.0673   -0.0564  0.09254   -0.0117    0.1313
## SalePriceMM          1.0000    0.1548  0.85661   -0.8532    0.0202
## SalePriceCH          0.1548    1.0000 -0.37714    0.0576   -0.7094
## PriceDiff            0.8566   -0.3771  1.00000   -0.8299    0.3894
## PctDiscMM           -0.8532    0.0576 -0.82995    1.0000    0.0158
## PctDiscCH            0.0202   -0.7094  0.38939    0.0158    1.0000
## ListPriceDiff        0.4441   -0.0621  0.44873   -0.1096   -0.0596
##           ListPriceDiff
## WeekofPurchase      0.0563
## PriceCH             -0.1646
## PriceMM              0.6682
## DiscCH              -0.0702
```

```
## DiscMM          -0.1005
## LoyalCH         0.0711
## SalePriceMM     0.4441
## SalePriceCH     -0.0621
## PriceDiff       0.4487
## PctDiscMM       -0.1096
## PctDiscCH       -0.0596
## ListPriceDiff   1.0000
```

```
corrplot(cor(data2), method = "color", type = "upper", tl.srt = 90)
```



3. Data Partition

```
set.seed(123)
trainindex <- createDataPartition(data1$Purchase, p=0.8, list= FALSE)
juice_train <- data1[trainindex, ]
juice_test <- data1[-trainindex, ]
dim(juice_train)
```

```
## [1] 800 18
```

```
dim(juice_test)
```

```
## [1] 200 18
```

4. SVM Model

a. Kernel = Linear

1. Cost = 0.01 (default)

```
svm1 <- svm(Purchase~., data=juice_train, kernel = "linear", cost=0.01)
summary(svm1)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = juice_train, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##
## Number of Support Vectors:  448
##
## ( 223 225 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

Training error rate:

```
train.pred = predict(svm1, juice_train)
x<-table(juice_train$Purchase, train.pred)
x
```

```
##      train.pred
##      CH  MM
## CH 436  52
## MM  82 230
```

```
a.linear<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
a.linear
```

```
## [1] 0.168
```

Test error rate:

```
test.pred = predict(svm1, juice_test)
x<-table(juice_test$Purchase, test.pred)
x
```

```
##      test.pred
##      CH  MM
##  CH 111  11
##  MM  22  56
```

```
b.linear<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
b.linear
```

```
## [1] 0.165
```

The training error rate is:

```
a.linear<-a.linear*100
capture.output(cat(a.linear, '%'))
```

```
## [1] "16.8 %"
```

The test error rate is:

```
b.linear<-b.linear*100
capture.output(cat(b.linear, '%'))
```

```
## [1] "16.5 %"
```

Tuning to select the best cost parameter

```
set.seed(123)
tune.out = tune(svm, Purchase ~ ., data = juice_train, kernel = "linear", ranges = list(cost = 10^seq(-1, 1, by = 0.25)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 0.562
##
## - best performance: 0.17
##
## - Detailed performance results:
##      cost error dispersion
## 1  0.0100 0.186      0.0419
## 2  0.0178 0.180      0.0417
## 3  0.0316 0.179      0.0354
## 4  0.0562 0.174      0.0291
```

```
## 5    0.1000 0.175    0.0333
## 6    0.1778 0.176    0.0330
## 7    0.3162 0.175    0.0312
## 8    0.5623 0.170    0.0271
## 9    1.0000 0.173    0.0262
## 10   1.7783 0.172    0.0269
## 11   3.1623 0.174    0.0224
## 12   5.6234 0.175    0.0212
## 13  10.0000 0.179    0.0250
```

Tuning shows that optimal cost is :

```
tune.out$best.parameters$cost
```

```
## [1] 0.562
```

2. SVM with best cost for kernel = linear

```
svm1_bestcost <- svm(Purchase~., data=juice_train, kernel = "linear", cost=tune.out$best.parameters$cost)
summary(svm1_bestcost)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = juice_train, kernel = "linear",
##      cost = tune.out$best.parameters$cost)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##           cost: 0.562
##
## Number of Support Vectors: 344
##
##   ( 171 173 )
##
##
## Number of Classes: 2
##
## Levels:
##   CH MM
```

Training error rate:

```
train.pred = predict(svm1_bestcost, juice_train)
x<-table(juice_train$Purchase, train.pred)
x
```



```
##      train.pred
##      CH  MM
##  CH 429  59
##  MM  74 238
```

```
a.linear.best<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
a.linear.best
```

```
## [1] 0.166
```

Test error rate:

```
test.pred = predict(svm1_bestcost, juice_test)
x<-table(juice_test$Purchase, test.pred)
x
```

```
##      test.pred
##      CH  MM
##  CH 109  13
##  MM  21  57
```

```
b.linear.best<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
b.linear.best
```

```
## [1] 0.17
```

The training error rate is:

```
a.linear.best<-a.linear.best*100
capture.output(cat(a.linear.best, '%'))
```

```
## [1] "16.6 %"
```

The test error rate is:

```
b.linear.best<-b.linear.best*100
capture.output(cat(b.linear.best, '%'))
```

```
## [1] "17 %"
```

The training error decreases to 16.6% but test error slightly increases from 16.5% to 17% by using best cost.

a. Kernel = radial

1. Cost = 0.01 (default)

```
set.seed(123)
svm2 <- svm(Purchase~., data=juice_train, kernel = "radial", cost=0.01)
summary(svm2)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = juice_train, kernel = "radial",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  0.01
##
## Number of Support Vectors:  626
##
##  ( 312 314 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

Training error rate:

```
train.pred = predict(svm2, juice_train)
x<-table(juice_train$Purchase, train.pred)
x
```

```
##      train.pred
##      CH  MM
## CH 488   0
## MM 312   0
```

```
a.radial<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
a.radial
```

```
## [1] 0.39
```

Test error rate:

```
test.pred = predict(svm2, juice_test)
x<-table(juice_test$Purchase, test.pred)
x
```

```
##      test.pred
##      CH  MM
## CH 122   0
## MM  78   0
```

```
b.radial<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
b.radial
```

```
## [1] 0.39
```

The training error rate is:

```
a.radial<-a.radial*100
capture.output(cat(a.radial, '%'))
```

```
## [1] "39 %"
```

The test error rate is:

```
b.radial<-b.radial*100
capture.output(cat(b.radial, '%'))
```

```
## [1] "39 %"
```

Tuning to select the best cost parameter

```
set.seed(123)
tune.out = tune(svm, Purchase ~ ., data = juice_train, kernel = "radial", ranges = list(cost = 10^seq(-1, 1, by = 0.25)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 0.562
##
## - best performance: 0.176
##
## - Detailed performance results:
##      cost error dispersion
## 1  0.0100 0.390    0.0642
## 2  0.0178 0.390    0.0642
## 3  0.0316 0.362    0.0757
## 4  0.0562 0.207    0.0401
## 5  0.1000 0.186    0.0458
## 6  0.1778 0.188    0.0391
## 7  0.3162 0.180    0.0364
## 8  0.5623 0.176    0.0365
## 9  1.0000 0.184    0.0404
```

```
## 10  1.7783 0.185      0.0394
## 11  3.1623 0.184      0.0363
## 12  5.6234 0.188      0.0358
## 13 10.0000 0.195      0.0355
```

Tuning shows that optimal cost is :

```
tune.out$best.parameters$cost
```

```
## [1] 0.562
```

2. SVM with best cost for kernel = radial

```
svm2_bestcost <- svm(Purchase~., data=juice_train, kernel = "radial", cost=tune.out$best.parameters$cost)
summary(svm2_bestcost)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = juice_train, kernel = "radial",
##      cost = tune.out$best.parameters$cost)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  0.562
##
## Number of Support Vectors:  412
##
##   ( 205 207 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

Training error rate:

```
train.pred = predict(svm2_bestcost, juice_train)
x<-table(juice_train$Purchase, train.pred)
x
```

```
##      train.pred
##      CH  MM
## CH 438  50
## MM  80 232
```

```
a.radial.best<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
a.radial.best
```

```
## [1] 0.163
```

Test error rate:

```
test.pred = predict(svm2_bestcost, juice_test)
x<-table(juice_test$Purchase, test.pred)
x
```

```
##      test.pred
##      CH  MM
## CH 113   9
## MM  22  56
```

```
b.radial.best<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
b.radial.best
```

```
## [1] 0.155
```

The training error rate is:

```
a.radial.best<-a.radial.best*100
capture.output(cat(a.radial.best, '%'))
```

```
## [1] "16.2 %"
```

The test error rate is:

```
b.radial.best<-b.radial.best*100
capture.output(cat(b.radial.best, '%'))
```

```
## [1] "15.5 %"
```

The training error decreases to 16.2% and test error decreases to 15.5% by using best cost.

a. Kernel = Poly
1. Cost = 0.01 (default)

```
svm3 <- svm(Purchase~., data=juice_train, kernel = "poly", cost=0.01, degree=2)
summary(svm3)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = juice_train, kernel = "poly",
##      cost = 0.01, degree = 2)
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel: polynomial
##      cost: 0.01
##      degree: 2
##      coef.0: 0
##
## Number of Support Vectors: 626
##
## ( 312 314 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

Training error rate:

```
train.pred = predict(svm3, juice_train)
x<-table(juice_train$Purchase, train.pred)
x
```

```
##      train.pred
##      CH  MM
## CH 488   0
## MM 312   0
```

```
a.poly<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
a.poly
```

```
## [1] 0.39
```

Test error rate:

```
test.pred = predict(svm3, juice_test)
x<-table(juice_test$Purchase, test.pred)
x
```

```
##      test.pred
##      CH  MM
## CH 122   0
## MM  78   0
```

```
b.poly<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
b.poly
```

```
## [1] 0.39
```

The training error rate is:

```
a.poly<-a.poly*100
capture.output(cat(a.poly, '%'))
```

```
## [1] "39 %"
```

The test error rate is:

```
b.poly<-b.poly*100
capture.output(cat(b.poly, '%'))
```

```
## [1] "39 %"
```

Tuning to select the best cost parameter

```
set.seed(123)
tune.out = tune(svm, Purchase ~ ., data = juice_train, kernel = "poly", degree= 2, ranges = list(cost =
  1, by = 0.25)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   5.62
##
## - best performance: 0.191
##
## - Detailed performance results:
##       cost error dispersion
## 1  0.0100 0.390    0.0642
## 2  0.0178 0.390    0.0642
## 3  0.0316 0.386    0.0639
## 4  0.0562 0.374    0.0630
## 5  0.1000 0.359    0.0483
## 6  0.1778 0.329    0.0574
## 7  0.3162 0.301    0.0649
## 8  0.5623 0.269    0.0563
## 9  1.0000 0.228    0.0489
```

```
## 10  1.7783 0.205      0.0392
## 11  3.1623 0.194      0.0461
## 12  5.6234 0.191      0.0453
## 13 10.0000 0.195      0.0413
```

Tuning shows that optimal cost is :

```
tune.out$best.parameters$cost
```

```
## [1] 5.62
```

2. SVM with best cost for kernel = poly

```
svm3_bestcost <- svm(Purchase~., data=juice_train, kernel = "poly", degree = 2, cost=tune.out$best.parameters$cost)
summary(svm3_bestcost)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = juice_train, kernel = "poly",
##      degree = 2, cost = tune.out$best.parameters$cost)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:  5.62
##    degree:  2
##    coef.0:  0
##
## Number of Support Vectors:  404
##
##   ( 197 207 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

Training error rate:

```
train.pred = predict(svm3_bestcost, juice_train)
x<-table(juice_train$Purchase, train.pred)
x
```

```
##      train.pred
##      CH  MM
## CH 453  35
## MM 103 209
```



```
a.poly.best<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
a.poly.best
```

```
## [1] 0.172
```

Test error rate:

```
test.pred = predict(svm3_bestcost, juice_test)
x<-table(juice_test$Purchase, test.pred)
x
```

```
##      test.pred
##      CH  MM
## CH 115   7
## MM  27  51
```

```
b.poly.best<-(x[2] + x[3])/(x[1] + x[2] + x[3] + x[4])
b.poly.best
```

```
## [1] 0.17
```

The training error rate is:

```
a.poly.best<-a.poly.best*100
capture.output(cat(a.poly.best, '%'))
```

```
## [1] "17.2 %"
```

The test error rate is:

```
b.poly.best<-b.poly.best*100
capture.output(cat(b.poly.best, '%'))
```

```
## [1] "17 %"
```

The training error decreases to 17.2% and test error decreases from 39% to 17% by using best cost.

```
rowLabels = c("Linear Kernel, Cost = 0.01",
              paste("Linear Kernel, with best cost"),
              "Radial Kernel, Cost = 0.01",
              paste("Radial Kernel, with best Cost"),
              "Polynomial Kernel, Degree = 2, Cost = 0.01",
              paste("Polynomial Kernel, Degree = 2, with best Cost"))

trainingErrorRate = c(a.linear,
                      a.linear.best,
                      a.radial,
                      a.radial.best,
```

```

a.poly,
a.poly.best)

testingErrorRate = c(b.linear,
b.linear.best,
b.radial,
b.radial.best,
b.poly,
b.poly.best)

df = data.frame(trainingErrorRate, testingErrorRate, row.names = rowLabels)
colnames(df) = c("Training Error Rate", "Testing Error Rate")
kable(df, format = 'markdown')

```

	Training Error Rate	Testing Error Rate
Linear Kernel, Cost = 0.01	16.8	16.5
Linear Kernel, with best cost	16.6	17.0
Radial Kernel, Cost = 0.01	39.0	39.0
Radial Kernel, with best Cost	16.2	15.5
Polynomial Kernel, Degree = 2, Cost = 0.01	39.0	39.0
Polynomial Kernel, Degree = 2, with best Cost	17.2	17.0

From the above table, overall, the Radial Kernel with Best Cost parameter of Cost = 0.562 seems to be producing minimum misclassification error on both training and testing data.

Explanation: From the Table, the SVM with a Radial Kernel and a cost of 0.562 seems to give the best results. It had the best train and test error rate compared to the other SVMs.

Looking at the other information in the table, the Radial and Poly SVMs for default cost = 0.01 had similar results. In other words, the SVM with a radial kernel of default cost = 0.01 had similar overall results to the SVM with a polynomial kernel of default cost = 0.01 and it should be noted that the SVMs with radial and polynomial kernels with default cost = 0.01 produced unreliable results without proper classification.